

# Data Mining Workflow

## Set Up the R Notebook for Analysis

```
# Load necessary packages
library('swat')

## SWAT 1.0.0

library('ggplot2')
library('reshape2')
options(cas.print.messages = FALSE)

# Data name
indata <- 'hmeq'

# Hostname, port, username, password
conn <- CAS(hostname, 8777, protocol = 'http')

## NOTE: Connecting to CAS and generating CAS action functions for loaded
##         action sets...

## NOTE: To generate the functions with signatures (for tab completion), set
##         options(cas.gen.function.sig=TRUE).

# Read in the dataset
castbl <- cas.read.csv(conn, paste0('http://support.sas.com/documentation',
                                     '/onlinedoc/viya/EXAMPLEDATASETS/hmeq.csv'))
```

## View Data

```
# Print the first few rows
head(castbl)
```

```
##   BAD LOAN MORTDUE  VALUE REASON   JOB  YOJ  DEROG DELINQ    CLAGE NINQ
## 1  1  1100    25860  39025 HomeImp Other 10.5    0    0  94.36667    1
## 2  1  1300    70053  68400 HomeImp Other  7.0    0    2 121.83333    0
## 3  1  1500   13500  16700 HomeImp Other  4.0    0    0 149.46667    1
## 4  1  1500     NaN    NaN           NaN   NaN   NaN     NaN   NaN
## 5  0  1700   97800 112000 HomeImp Office 3.0    0    0  93.33333    0
## 6  1  1700   30548  40320 HomeImp Other  9.0    0    0 101.46600    1
##   CLNO  DEBTINC
## 1    9     NaN
## 2   14     NaN
## 3   10     NaN
## 4  NaN     NaN
## 5   14     NaN
## 6    8 37.11361
```

## Get Summary Statistics


```
# Use summary function to get variable summary
summary(castbl)
```

```
##          BAD          LOAN          MORTDUE
## Min.      :0.0000   Min.    : 1100   Min.      :2063
## 1st Qu.:0.0000   1st Qu.:11100   1st Qu.:46268
## Median :0.0000   Median :16300   Median :65019
## Mean     :0.1995   Mean    :18608   Mean     :73760.8171995589
## 3rd Qu.:0.0000   3rd Qu.:23300   3rd Qu.:91491
## Max.      :1.0000   Max.     :89900   Max.     :399550
##                                     NA's      :518
##          VALUE          REASON          JOB
## Min.      :8000          DebtCon:3928   Mgr       : 767
## 1st Qu.:66069          HomeImp:1780   Office    : 948
## Median :89235.5        NA's       : 252   Other     :2388
## Mean     :101776.04874145          ProfExe:1276
## 3rd Qu.:119831.5          Sales     : 109
## Max.      :855909          Self      : 193
## NA's      :112          NA's       : 279
##          YOJ          DEROG
## Min.      :0          Min.      :0
## 1st Qu.:3          1st Qu.:0
## Median :7          Median :0
## Mean     :8.92226813590449   Mean    :0.254569687738
## 3rd Qu.:13          3rd Qu.:0
## Max.      :41          Max.      :10
## NA's      :515        NA's      :708
##          DELINQ          CLAGE
## Min.      :0          Min.      :0
## 1st Qu.:0          1st Qu.:115.103196832924
## Median :0          Median :173.466666666667
## Mean     :0.44944237918215   Mean    :179.766275186577
## 3rd Qu.:0          3rd Qu.:231.574833599946
## Max.      :15          Max.      :1168.23356094464
## NA's      :580        NA's      :308
##          NINQ          CLNO
## Min.      :0          Min.      :0
## 1st Qu.:0          1st Qu.:15
## Median :1          Median :20
## Mean     :1.18605504587155   Mean    :21.2960962007668
## 3rd Qu.:2          3rd Qu.:26
## Max.      :17          Max.      :71
## NA's      :510        NA's      :222
##          DEBTINC
## Min.      :0.52449921542988
## 1st Qu.:29.1400313718617
## Median :34.818261818587
## Mean     :33.7799153487192
## 3rd Qu.:39.0031406283719
## Max.      :203.312148691165
## NA's      :1267
```

## Visualize Numeric Variables

```
# Bring data locally
df <- to.casDataFrame(castbl, obs = nrow(castbl))

# Use reshape2's melt to help with data formatting
d <- melt(df[sapply(df, is.numeric)], id.vars=NULL)
ggplot(d, aes(x = value)) +
  facet_wrap(~variable, scales = 'free_x') +
  geom_histogram(fill = 'blue', bins = 25)
```




## Check for Missingness

```
# Check for missing values
tbl <- cas.simple.distinct(castbl)$Distinct[,c('Column', 'NMiss')]
tbl
```

##	Column	NMiss
## 1	BAD	0
## 2	LOAN	0
## 3	MORTDUE	518
## 4	VALUE	112
## 5	REASON	252
## 6	JOB	279
## 7	YOJ	515

```
## 8      DEROG      708
## 9      DELINQ      580
## 10     CLAGE      308
## 11     NINQ       510
## 12     CLNO       222
## 13 DEBTINC      1267
```

```
# Visualize the missing data
tbl$PctMiss <- tbl$NMiss/nrow(castbl)
ggplot(tbl, aes(Column, PctMiss)) +
  geom_col(fill = 'blue') +
  ggtitle('Pct Missing Values') +
  theme(plot.title = element_text(hjust = 0.5))
```



## Impute Missing Values

```
# Impute missing values
cas.dataPreprocess.impute(castbl,
  methodContinuous = 'MEDIAN',
  methodNominal    = 'MODE',
  inputs           = colnames(castbl)[-1],
  copyAllVars      = TRUE,
  casOut           = list(name = indata, replace = TRUE)
)
```

```
## $ImputeInfo
```

```

##      Variable ImputeTech      ResultVar      N NMiss ImputedValueContinuous
## 1      LOAN      Median      IMP_LOAN 5960      0      16300.00000
## 2     MORTDUE      Median IMP_MORTDUE 5442    518      65019.00000
## 3      VALUE      Median      IMP_VALUE 5848    112      89235.50000
## 4     REASON      Mode      IMP_REASON 5708    252           NaN
## 5       JOB      Mode      IMP_JOB 5681    279           NaN
## 6      YOJ      Median      IMP_YOJ 5445    515       7.00000
## 7     DEROG      Median      IMP_DEROG 5252    708       0.00000
## 8     DELINQ      Median      IMP_DELINQ 5380    580       0.00000
## 9      CLAGE      Median      IMP_CLAGE 5652    308     173.46667
## 10     NINQ      Median      IMP_NINQ 5450    510       1.00000
## 11     CLNO      Median      IMP_CLNO 5738    222     20.00000
## 12    DEBTINC      Median      IMP_DEBTINC 4693  1267     34.81826
##      ImputedValueNominal
## 1
## 2
## 3
## 4      DebtCon
## 5      Other
## 6
## 7
## 8
## 9
## 10
## 11
## 12
##
## $OutputCasTables
##      casLib Name Rows Columns
## 1 CASUSERHDFS(jelueb) hmeq 5960      25

```

## Split the Data into Training and Validation

```

# Load the sampling actionset
loadActionSet(conn, 'sampling')

# Partition the data
cas.sampling.srs(conn,
  table = indata,
  sampct = 30,
  partind = TRUE,
  output = list(casOut = list(name = indata, replace = T), copyVars = 'ALL')
)

```

```

# Load the fedsq actionset
loadActionSet(conn, 'fedsq')

# Make sure the partition worked correctly using SQL
cas.fedsq.execDirect(conn, query = paste0("
  SELECT
    CASE WHEN _PartInd_ = 0 THEN 'Training' ELSE 'Validation' END AS name,
    _PartInd_,
    COUNT(*) AS obs

```

```
FROM ", indata, "
GROUP BY
    CASE WHEN _PartInd_ = 0 THEN 'Training' ELSE 'Validation' END,
    _PartInd_;
"))$`Result Set`
```

```
##          NAME _PartInd_  OBS
## 1   Training          0 4172
## 2 Validation          1 1788
```

## Variable Shortcuts

Note: I do not want to hard code any of my variable names.

```
# Get variable info and types
colinfo <- head(cas.table.columnInfo(conn, table = indata)$ColumnInfo, -1)

# My target variable is the first column
target <- colinfo$Column[1]

# For models that can inherently handle missing values (ex: Decision Tree)
inputs <- colinfo$Column[-1]
nominals <- c(target, subset(colinfo, Type == 'varchar')$Column)

# For models that cannot handle missing values (ex: Neural Network)
imp.inputs <- grep('IMP_', inputs, value = T)
imp.nominals <- c(target, grep('IMP_', nominals, value = T))
```

## Model Building

### Decision Tree

```
# Load the decision tree actionset
loadActionSet(conn, 'decisionTree')

# Train the decision tree model
cas.decisionTree.dtreeTrain(conn,
    table = list(name = indata, where = '_PartInd_ = 0'),
    target = target,
    inputs = inputs,
    nominals = nominals,
    varImp = TRUE,
    casOut = list(name = 'dt_model', replace = TRUE)
)
```

```
## $DTreeVarImpInfo
##   Variable Importance      Std Count
## 1  DEBTINC 438.877290 181.520795      2
## 2   DELINQ  50.426019  24.948820      2
## 3    DEROG  22.090378   2.396851      2
## 4    VALUE   8.621421   0.000000      1
```

```
## 5      NINQ      1.716695      0.000000      1
##
## $ModelInfo
##              Descr      Value
## 1      Number of Tree Nodes      17.00000
## 2      Max Number of Branches      2.00000
## 3      Number of Levels      6.00000
## 4      Number of Leaves      9.00000
## 5      Number of Bins      20.00000
## 6      Minimum Size of Leaves      5.00000
## 7      Maximum Size of Leaves      3203.00000
## 8      Number of Variables      24.00000
## 9      Confidence Level for Pruning      0.25000
## 10     Number of Observations Used      4172.00000
## 11     Misclassification Error (%)      13.68648
##
## $OutputCasTables
##              casLib      Name Rows Columns
## 1 CASUSERHDFS(jelueb) dt_model      17      27
```

## Random Forest

```
# Train the random forest model
cas.decisionTree.forestTrain(conn,
  table      = list(name = indata, where = '_PartInd_ = 0'),
  target     = target,
  inputs     = inputs,
  nominals   = nominals,
  casOut     = list(name = 'rf_model', replace = TRUE)
)
```

```
## $ModelInfo
##              Descr      Value
## 1      Number of Trees      50.00000
## 2      Number of Selected Variables (M)      5.00000
## 3      Random Number Seed      0.00000
## 4      Bootstrap Percentage (%)      63.21206
## 5      Number of Bins      20.00000
## 6      Number of Variables      24.00000
## 7      Confidence Level for Pruning      0.25000
## 8      Max Number of Tree Nodes      29.00000
## 9      Min Number of Tree Nodes      11.00000
## 10     Max Number of Branches      2.00000
## 11     Min Number of Branches      2.00000
## 12     Max Number of Levels      6.00000
## 13     Min Number of Levels      6.00000
## 14     Max Number of Leaves      15.00000
## 15     Min Number of Leaves      6.00000
## 16     Maximum Size of Leaves      2590.00000
## 17     Minimum Size of Leaves      5.00000
## 18     Out-of-Bag MCR (%)      NaN
##
## $OutputCasTables
```

```
##          casLib      Name Rows Columns
## 1 CASUSERHDFS(jelueb) rf_model  742      41
```

## Gradient Boosting

```
# Train the gradient boosting model
cas.decisionTree.gbtTreeTrain(conn,
  table   = list(name = indata, where = '_PartInd_ = 0'),
  target  = target,
  inputs  = inputs,
  nominals = nominals,
  casOut  = list(name = 'gbt_model', replace = TRUE)
)
```

```
## $ModelInfo
##          Descr  Value
## 1      Number of Trees  50.0
## 2      Distribution     2.0
## 3      Learning Rate    0.1
## 4      Subsampling Rate  0.5
## 5 Number of Selected Variables (M) 24.0
## 6      Number of Bins    20.0
## 7      Number of Variables 24.0
## 8      Max Number of Tree Nodes 61.0
## 9      Min Number of Tree Nodes 29.0
## 10     Max Number of Branches   2.0
## 11     Min Number of Branches   2.0
## 12     Max Number of Levels     6.0
## 13     Min Number of Levels     6.0
## 14     Max Number of Leaves    31.0
## 15     Min Number of Leaves    15.0
## 16     Maximum Size of Leaves 1492.0
## 17     Minimum Size of Leaves   5.0
## 18     Random Number Seed     0.0
##
## $OutputCasTables
##          casLib      Name Rows Columns
## 1 CASUSERHDFS(jelueb) gbt_model 2492      31
```

## Neural Network

```
# Load the neuralNet actionset
loadActionSet(conn, 'neuralNet')

# Build a neural network model
cas.neuralNet.annTrain(conn,
  table   = list(name = indata, where = '_PartInd_ = 0'),
  target  = target,
  inputs  = imp.inputs,
  nominals = imp.nominals,
  casOut  = list(name = 'nn_model', replace = TRUE)
)
```



```
## $ConvergenceStatus
##                                     Reason
## 1 The optimization exited on maximum iterations.
##
## $ModelInfo
##           Descr      Value
## 1           Model  Neural Net
## 2 Number of Observations Used      4172
## 3 Number of Observations Read      4172
## 4   Target/Response Variable      BAD
## 5       Number of Nodes           20
## 6   Number of Input Nodes          18
## 7   Number of Output Nodes          2
## 8   Number of Hidden Nodes          0
## 9 Number of Weight Parameters       18
## 10 Number of Bias Parameters         2
## 11           Architecture      GLIM
## 12   Number of Neural Nets           1
## 13           Objective Value 1.5329196041
##
## $OptIterHistory
##   Progress Objective      Loss
## 1         1  4.184992  4.184992
## 2         2  2.456393  2.456393
## 3         3  1.658346  1.658346
## 4         4  1.584687  1.584687
## 5         5  1.548499  1.548499
## 6         6  1.540727  1.540727
## 7         7  1.535718  1.535718
## 8         8  1.534470  1.534470
## 9         9  1.533433  1.533433
## 10        10  1.532920  1.532920
##
## $OutputCasTables
##           casLib      Name Rows Columns
## 1 CASUSERHDFS(jelueb) nn_model    20     15
```

## Score the Models

```
# Score the models
models <- c('dt','rf','gbt','nn')
scores <- c(cas.decisionTree.dtreeScore, cas.decisionTree.forestScore,
            cas.decisionTree.gbtTreeScore, cas.neuralNet.annScore)
names(scores) <- models

# Function to help automate prediction process on new data
score.params <- function(model){return(list(
  object      = defCasTable(conn, indata),
  modelTable  = list(name = paste0(model, '_model')),
  copyVars    = list(target, '_PartInd_'),
  assessorRow = TRUE,
  casOut      = list(name = paste0(model, '_scored'), replace = T)
))}
```

```
lapply(models, function(x) {do.call(scores[[x]], score.params(x))})
```

## Compare Confusion Matrix

```
# Load the percentile actionset for scoring
loadActionSet(conn, 'percentile')

# Useful function for model assessment
assess.model <- function(model){
  cas.percentile.assess(conn,
    table = list(name = paste0(model, '_scored'),
                  where = '_PartInd_ = 1'),
    inputs = paste0('_', model, '_P_', 1),
    response = target,
    event = '1')
}

model.names <- c('Decision Tree', 'Random Forest',
                 'Gradient Boosting', 'Neural Network')
roc.df <- data.frame()
for (i in 1:length(models)){
  tmp <- (assess.model(models[i]))$ROCInfo
  tmp$Model <- model.names[i]
  roc.df <- rbind(roc.df, tmp)
}

# Manipulate the dataframe
compare <- subset(roc.df, CutOff == 0.5)
rownames(compare) <- NULL
compare[,c('Model', 'TP', 'FP', 'FN', 'TN')]
```

```
##           Model TP  FP  FN  TN
## 1  Decision Tree 244 152 103 1289
## 2   Random Forest   35   1 312 1440
## 3 Gradient Boosting 224  58 123 1383
## 4   Neural Network 106  50 241 1391
```

## Compare Misclassification

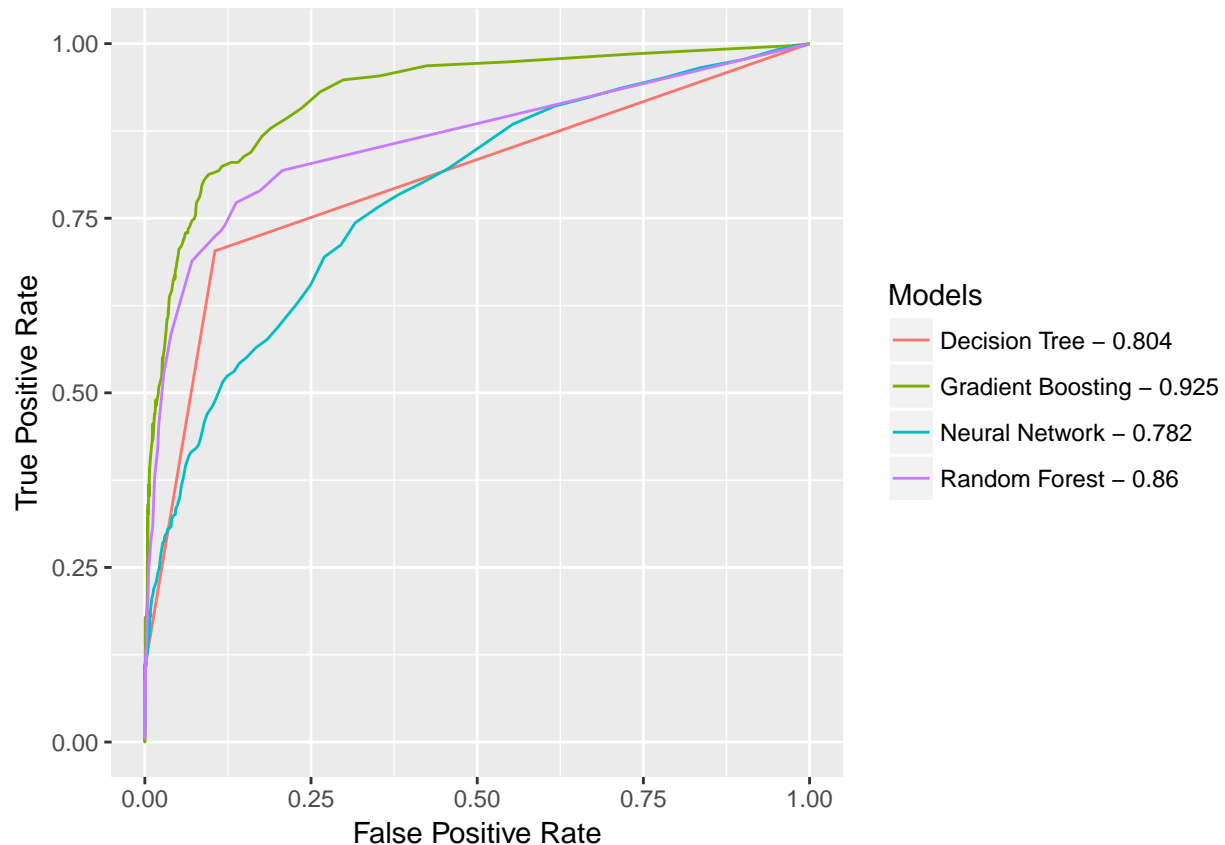
```
# Build a dataframe to compare the misclassification rates
compare$Misclassification <- 1 - compare$ACC
miss <- compare[order(compare$Misclassification), c('Model', 'Misclassification')]
rownames(miss) <- NULL
miss
```

```
##           Model Misclassification
## 1 Gradient Boosting           0.1012304
## 2   Decision Tree           0.1426174
## 3   Neural Network           0.1627517
## 4   Random Forest           0.1750559
```

## Compare ROC Curve

```
# Add a new column to be used as the ROC curve label
roc.df$Models <- paste(roc.df$Model, round(roc.df$C, 3), sep = ' - ')

# Create the ROC curve
ggplot(data = roc.df[c('FPR', 'Sensitivity', 'Models')],
       aes(x = as.numeric(FPR), y = as.numeric(Sensitivity), colour = Models)) +
  geom_line() +
  labs(x = 'False Positive Rate', y = 'True Positive Rate')
```



## Save the CAS Gradient Boosting Model

```
# Save the champion model for later use
cas.table.save(conn, table = list(name = 'gbt_model'), name = 'gbt_model', replace = T)

## $caslib
## [1] "CASUSERHDFS(jelueb)"
##
## $name
## [1] "gbt_model.sashdat"
```

## End the Session

```
# End the session  
cas.session.endSession(conn)
```