

Data Mining Workflow

Set Up the R Notebook for Analysis

```
# Load necessary packages
library('swat')

## SWAT 1.0.0

library('ggplot2')
library('reshape2')
options(cas.print.messages = FALSE)

# Hostname, port, username, password
conn <- CAS(hostname, port, username, password)

## NOTE: Connecting to CAS and generating CAS action functions for loaded
##        action sets...

## NOTE: To generate the functions with signatures (for tab completion), set
##        options(cas.gen.function.sig=TRUE).

# Change the active caslib to public
cas.sessionProp.setSessOpt(conn, caslib = 'public')
```

View Data

```
# Create a CAS table for the prepped data set
castbl <- defCasTable(conn, 'hmeq_prepped')

# Print the first few rows
head(castbl)
```

	BAD	LOAN	MORTDUE	VALUE	REASON	JOB	YOJ	DEROG	DELINQ	CLAGE	NINQ
## 1	1	1100	25860	39025	HomeImp	Other	10.5	0	0	94.36667	1
## 2	1	1500	NaN	NaN			NaN	NaN	NaN	NaN	NaN
## 3	1	1800	48649	57037	HomeImp	Other	5.0	3	2	77.10000	1
## 4	1	2000	NaN	62250	HomeImp	Sales	16.0	0	0	115.80000	0
## 5	1	2000	45000	55000	HomeImp	Other	3.0	0	0	86.06667	2
## 6	1	2200	24280	34687	HomeImp	Other	NaN	0	1	300.86667	0
	CLNO	DEBTINC	IMP_CLAGE	IMP_CLNO	IMP_DEBTINC	IMP_DELINQ	IMP_DEROG				
## 1	9	NaN	94.36667	9	34.81826	0	0				
## 2	NaN	NaN	173.46667	20	34.81826	0	0				
## 3	17	NaN	77.10000	17	34.81826	2	3				
## 4	13	NaN	115.80000	13	34.81826	0	0				
## 5	25	NaN	86.06667	25	34.81826	0	0				
## 6	8	NaN	300.86667	8	34.81826	1	0				
	IMP_LOAN	IMP_MORTDUE	IMP_MORTPAID	IMP_NINQ	IMP_VALUE	IMP_YOJ	IMP_JOB				
## 1	1100	25860	13165	1	39025.0	10.5	Other				
## 2	1500	65019	26623	1	89235.5	7.0	Other				
## 3	1800	48649	8388	1	57037.0	5.0	Other				

```
## 4      2000      65019      26623      0  62250.0    16.0  Sales
## 5      2000      45000      10000      2  55000.0     3.0  Other
## 6      2200      24280      10407      0  34687.0     7.0  Other
##   IMP_REASON MORTPAID _PartInd_
## 1   HomeImp    13165         1
## 2   DebtCon      NaN         0
## 3   HomeImp     8388         1
## 4   HomeImp      NaN         0
## 5   HomeImp    10000         0
## 6   HomeImp    10407         0
```

Variable Shortcuts

Note: I do not want to hard code any of my variable names.

```
# Get variable info and types
colinfo <- head(cas.table.columnInfo(conn, table = 'hmeq_prepped')$ColumnInfo, -1)

# My target variable is the first column
target <- colinfo$Column[1]
vars <- colinfo$Column[-1]
noms <- c(target, subset(colinfo, Type == 'varchar')$Column)

# For models that can inherently handle missing values (ex: Decision Tree)
inputs <- grep('IMP_', vars, value = TRUE, invert = TRUE)
nominals <- grep('IMP_', noms, value = TRUE, invert = TRUE)

# For models that cannot handle missing values (ex: Neural Network)
imp.inputs <- grep('IMP_', vars, value = TRUE)
imp.nominals <- c(target, grep('IMP_', noms, value = TRUE))
```

Model Building

Decision Tree

```
# Load the decision tree actionset
loadActionSet(conn, 'decisionTree')

# Train the decision tree model
cas.decisionTree.dtreeTrain(conn,
  table = list(name = 'hmeq_prepped', where = '_PartInd_ = 0'),
  target = target,
  inputs = inputs,
  nominals = nominals,
  varImp = TRUE,
  casOut = list(name = 'dt_model', replace = TRUE)
)

## $DTreeVarImpInfo
##   Variable Importance      Std Count
## 1  DEBTINC 442.166677 183.520139      2
## 2   DELINQ  47.167293   0.000000      1
```

```
## 3      DEROG  11.843676   2.650169     2
## 4      VALUE   8.545419   0.000000     1
## 5     MORTDUE   1.819619   0.000000     1
##
## $ModelInfo
##              Descr      Value
## 1      Number of Tree Nodes  15.00000
## 2      Max Number of Branches  2.00000
## 3      Number of Levels      6.00000
## 4      Number of Leaves      8.00000
## 5      Number of Bins       20.00000
## 6      Minimum Size of Leaves  5.00000
## 7      Maximum Size of Leaves 3214.00000
## 8      Number of Variables   13.00000
## 9      Confidence Level for Pruning  0.25000
## 10     Number of Observations Used 4172.00000
## 11     Misclassification Error (%)  13.75839
##
## $OutputCasTables
##   casLib      Name Rows Columns
## 1 Public dt_model   15      27
```

Random Forest

```
# Train the random forest model
cas.decisionTree.forestTrain(conn,
  table   = list(name = 'hmeq_prepped', where = '_PartInd_ = 0'),
  target  = target,
  inputs  = inputs,
  nominals = nominals,
  casOut  = list(name = 'rf_model', replace = TRUE)
)
```

```
## $ModelInfo
##              Descr      Value
## 1      Number of Trees      50.00000
## 2     Number of Selected Variables (M)  4.00000
## 3      Random Number Seed      0.00000
## 4      Bootstrap Percentage (%)  63.21206
## 5      Number of Bins       20.00000
## 6      Number of Variables   13.00000
## 7      Confidence Level for Pruning  0.25000
## 8      Max Number of Tree Nodes  33.00000
## 9      Min Number of Tree Nodes  11.00000
## 10     Max Number of Branches    2.00000
## 11     Min Number of Branches    2.00000
## 12     Max Number of Levels      6.00000
## 13     Min Number of Levels      6.00000
## 14     Max Number of Leaves     17.00000
## 15     Min Number of Leaves      6.00000
## 16     Maximum Size of Leaves 2580.00000
## 17     Minimum Size of Leaves    5.00000
## 18     Out-of-Bag MCR (%)       NaN
```

```
##
## $OutputCasTables
##   casLib      Name Rows Columns
## 1 Public rf_model  852      41
```

Gradient Boosting

```
# Train the gradient boosting model
cas.decisionTree.gbtTreeTrain(conn,
  table   = list(name = 'hmeq_prepped', where = '_PartInd_ = 0'),
  target  = target,
  inputs  = inputs,
  nominals = nominals,
  casOut  = list(name = 'gbt_model', replace = TRUE)
)
```

```
## $ModelInfo
##
##           Descr  Value
## 1      Number of Trees  50.0
## 2      Distribution     2.0
## 3      Learning Rate    0.1
## 4      Subsampling Rate  0.5
## 5 Number of Selected Variables (M) 13.0
## 6      Number of Bins   20.0
## 7      Number of Variables 13.0
## 8      Max Number of Tree Nodes 61.0
## 9      Min Number of Tree Nodes 31.0
## 10     Max Number of Branches   2.0
## 11     Min Number of Branches   2.0
## 12     Max Number of Levels     6.0
## 13     Min Number of Levels     6.0
## 14     Max Number of Leaves    31.0
## 15     Min Number of Leaves    16.0
## 16     Maximum Size of Leaves 1736.0
## 17     Minimum Size of Leaves   5.0
## 18     Random Number Seed     0.0
##
## $OutputCasTables
##   casLib      Name Rows Columns
## 1 Public gbt_model 2470      31
```

Neural Network

```
# Load the neuralNet actionset
loadActionSet(conn, 'neuralNet')

# Build a neural network model
cas.neuralNet.annTrain(conn,
  table   = list(name = 'hmeq_prepped', where = '_PartInd_ = 0'),
  target  = target,
  inputs  = imp.inputs,
  nominals = imp.nominals,
```

```
casOut = list(name = 'nn_model', replace = TRUE)
)
```

```
## $ConvergenceStatus
##                                     Reason
## 1 The optimization exited on maximum iterations.
##
## $ModelInfo
##           Descr      Value
## 1           Model  Neural Net
## 2 Number of Observations Used      4172
## 3 Number of Observations Read      4172
## 4   Target/Response Variable      BAD
## 5           Number of Nodes        21
## 6       Number of Input Nodes        19
## 7       Number of Output Nodes         2
## 8       Number of Hidden Nodes         0
## 9 Number of Weight Parameters        19
## 10 Number of Bias Parameters          2
## 11           Architecture      GLIM
## 12       Number of Neural Nets         1
## 13           Objective Value 1.5687790628
##
## $OptIterHistory
##   Progress Objective      Loss
## 1         1  4.568050  4.568050
## 2         2  2.879156  2.879156
## 3         3  1.748308  1.748308
## 4         4  1.660549  1.660549
## 5         5  1.613690  1.613690
## 6         6  1.595178  1.595178
## 7         7  1.578170  1.578170
## 8         8  1.573189  1.573189
## 9         9  1.569972  1.569972
## 10        10  1.568779  1.568779
##
## $OutputCasTables
##   casLib      Name Rows Columns
## 1 Public nn_model   21      15
```

Score the Models

```
# Score the models
models <- c('dt','rf','gbt','nn')
scores <- c(cas.decisionTree.dtreeScore, cas.decisionTree.forestScore,
            cas.decisionTree.gbtTreeScore, cas.neuralNet.annScore)
names(scores) <- models

# Function to help automate prediction process on new data
score.params <- function(model){return(list(
  object      = defCasTable(conn, 'hmeq_prepped'),
  modelTable  = list(name = paste0(model, '_model')),
```

```

copyVars      = list(target, '_PartInd_'),
assessonerow = TRUE,
casOut        = list(name = paste0(model, '_scored'), replace = T)
))}
lapply(models, function(x) {do.call(scores[[x]], score.params(x))})

```

Compare Confusion Matrix

```

# Load the percentile actionset for scoring
loadActionSet(conn, 'percentile')

# Useful function for model assessment
assess.model <- function(model){
  cas.percentile.assess(conn,
    table      = list(name = paste0(model, '_scored'),
                      where = '_PartInd_ = 1'),
    inputs     = paste0('_', model, '_P_', 1),
    response   = target,
    event      = '1')
}

model.names <- c('Decision Tree', 'Random Forest',
                 'Gradient Boosting', 'Neural Network')
roc.df <- data.frame()
for (i in 1:length(models)){
  tmp <- (assess.model(models[i]))$ROCInfo
  tmp$Model <- model.names[i]
  roc.df <- rbind(roc.df, tmp)
}

# Manipulate the dataframe
compare <- subset(roc.df, CutOff == 0.5)
rownames(compare) <- NULL
compare[,c('Model', 'TP', 'FP', 'FN', 'TN')]

```

```

##           Model TP  FP  FN  TN
## 1  Decision Tree 244 163  89 1292
## 2   Random Forest 125  13 208 1442
## 3 Gradient Boosting 234  63  99 1392
## 4   Neural Network 116  51 217 1404

```

Compare Misclassification

```

# Build a dataframe to compare the misclassification rates
compare$Misclassification <- 1 - compare$ACC
miss <- compare[order(compare$Misclassification), c('Model', 'Misclassification')]
rownames(miss) <- NULL
miss

```

```

##           Model Misclassification
## 1 Gradient Boosting          0.09060403

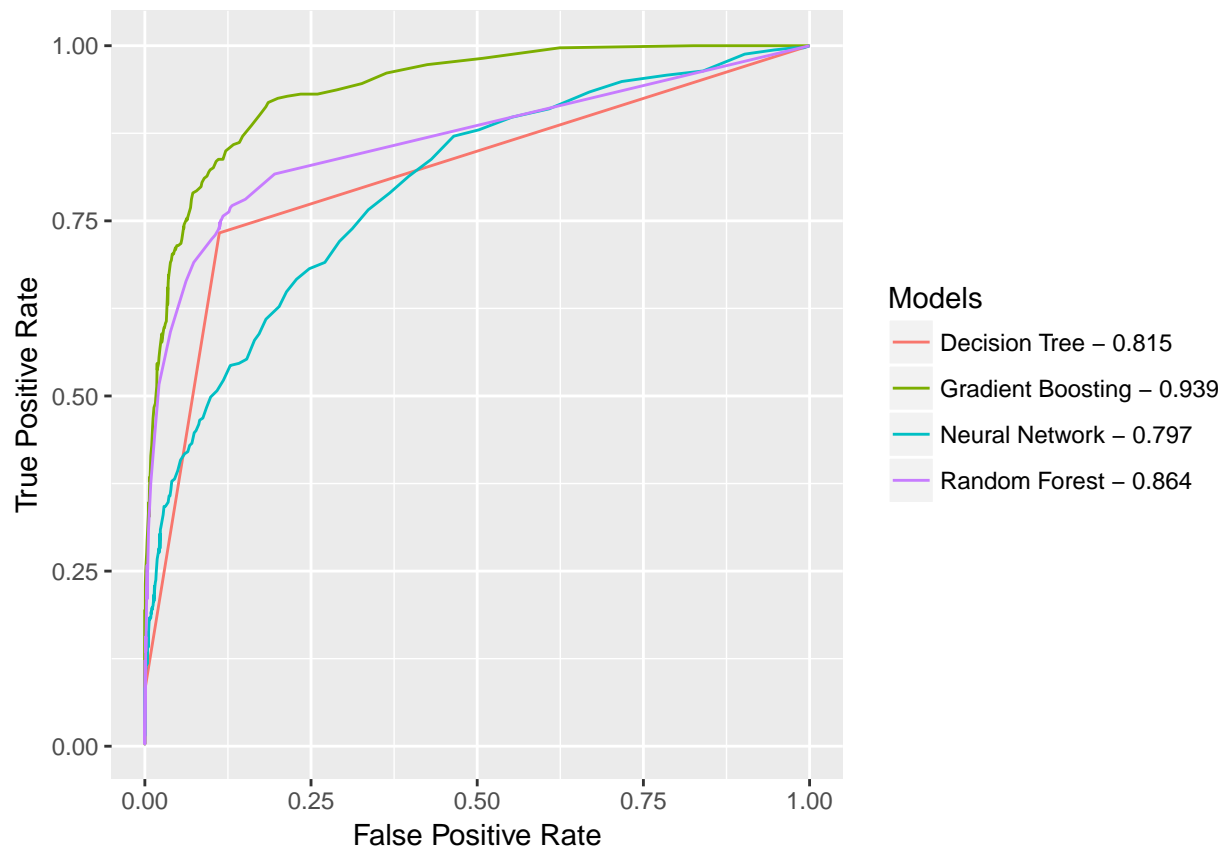
```

```
## 2    Random Forest      0.12360179
## 3    Decision Tree      0.14093960
## 4    Neural Network     0.14988814
```

Compare ROC Curve

```
# Add a new column to be used as the ROC curve label
roc.df$Models <- paste(roc.df$Model, round(roc.df$C, 3), sep = ' - ')

# Create the ROC curve
ggplot(data = roc.df[c('FPR', 'Sensitivity', 'Models')],
       aes(x = as.numeric(FPR), y = as.numeric(Sensitivity), colour = Models)) +
  geom_line() +
  labs(x = 'False Positive Rate', y = 'True Positive Rate')
```



Compare XGBoost Model

```
library('xgboost')
suppressPackageStartupMessages(library('caret'))

# Bring data to R client
df <- to.casDataFrame(castbl, obs = nrow(castbl))
df <- df[,c(target, inputs, '_PartInd_')]
```

```

# Create dummy variables through one-hot encoding
df.dum <- df[,nominals[-1]]
dummies <- dummyVars('~ .', data = df.dum)
df.ohe <- as.data.frame(predict(dummies, newdata = df))
df.all.combined <- cbind(df[, -c(which(colnames(df) %in% nominals[-1]))], df.ohe)

# Split into training and validation
train <- df.all.combined[df.all.combined['_PartInd_'] == 0,]
valid <- df.all.combined[df.all.combined['_PartInd_'] == 1,]

# Train the XGBoost model
bst <- xgboost(
  data = data.matrix(train[,-1]),
  label = data.matrix(train[,1]),
  objective = "binary:logistic",
  nround = 50,
  eta = 0.1,
  subsample = 0.5,
  colsample_bytree = 0.5
)

## [1] train-error:0.107143
## [2] train-error:0.100192
## [3] train-error:0.098754
## [4] train-error:0.100671
## [5] train-error:0.095638
## [6] train-error:0.093720
## [7] train-error:0.089406
## [8] train-error:0.090364
## [9] train-error:0.087248
## [10] train-error:0.086769
## [11] train-error:0.086050
## [12] train-error:0.082694
## [13] train-error:0.083174
## [14] train-error:0.080297
## [15] train-error:0.079338
## [16] train-error:0.079818
## [17] train-error:0.079578
## [18] train-error:0.077900
## [19] train-error:0.074065
## [20] train-error:0.074305
## [21] train-error:0.073826
## [22] train-error:0.074065
## [23] train-error:0.073346
## [24] train-error:0.072627
## [25] train-error:0.072148
## [26] train-error:0.071668
## [27] train-error:0.070470
## [28] train-error:0.069271
## [29] train-error:0.068792
## [30] train-error:0.068073
## [31] train-error:0.068552
## [32] train-error:0.066395

```



```
## [33] train-error:0.065436
## [34] train-error:0.064477
## [35] train-error:0.063279
## [36] train-error:0.061361
## [37] train-error:0.061361
## [38] train-error:0.060163
## [39] train-error:0.059923
## [40] train-error:0.059444
## [41] train-error:0.058965
## [42] train-error:0.058245
## [43] train-error:0.056807
## [44] train-error:0.057287
## [45] train-error:0.057287
## [46] train-error:0.056568
## [47] train-error:0.054890
## [48] train-error:0.055129
## [49] train-error:0.054410
## [50] train-error:0.053931
```

Score and Assess XGBoost on Validation Data

```
# Create a dataframe with the misclassification rate for XGBoost
pred <- as.numeric(predict(bst, data.matrix(valid[,-1]), missing = 'NAN') > 0.5)
Misclassification <- mean(as.numeric(pred > 0.5) != valid[,1])
xgb <- data.frame(cbind(Model = 'R - XGBoost', Misclassification))
xgb
```

```
##           Model Misclassification
## 1 R - XGBoost 0.0911633109619687
```

Final Assessment with CAS and R Models

```
# Combine the assessments and order by most accurate on validation data
err <- data.frame(rbind(miss, xgb))
err[,-1] <- round(as.numeric(as.character(err[,-1])),7)
err <- err[order(err[,-1]),]
rownames(err) <- NULL
err
```

```
##           Model Misclassification
## 1 Gradient Boosting      0.0906040
## 2      R - XGBoost      0.0911633
## 3      Random Forest    0.1236018
## 4      Decision Tree    0.1409396
## 5      Neural Network    0.1498881
```

Save the CAS Gradient Boosting Model

```
# Save the champion model to disk for later use
cas.table.save(conn, table = list(name = 'gbt_model'), name = 'Jesse_SAS_gbt', replace = T)
```

```
## $caslib
## [1] "Public"
##
## $name
## [1] "Jesse_SAS_gbt.sashdat"
# Promote the champion model to public memory to share with team
cas.table.promote(conn, name = 'gbt_model', target = 'Jesse_SAS_gbt', targetLib = 'public')

## list()
# Save the challenger (XGBoost) model for later use
xgb.save(bst, "Jesse_R_xgb.model")

## [1] TRUE
```

End the Session

```
# End the session
cas.session.endSession(conn)
```